

About me

- David Kröll, IT student
- Docker user since Jun `17

<https://davidkroell.github.io>



Preface

- Container technology highly used in cloud
- Must-have for Ops in 2018
- Presentation and snippets
 - <https://github.com/davidkroell/docker-deepdive>

Topics

- General
- Difference to VMs
- Image vs. Container
- Registry
- Dockerfile
- docker-compose

Topics

- Storage
- Networking
- Swarm
- Monitoring
- Outlook

General

- Collection of tools
 - runc, containerd
- Versions
 - CE, EE
- Builds
 - edge, stable
- written in Golang

General

- Save costs in application deployment
 - Same motivation as VMs
- Isolation
- „Runs on my machine“
 - Packaging of apps
 - Build, Ship, Run



General

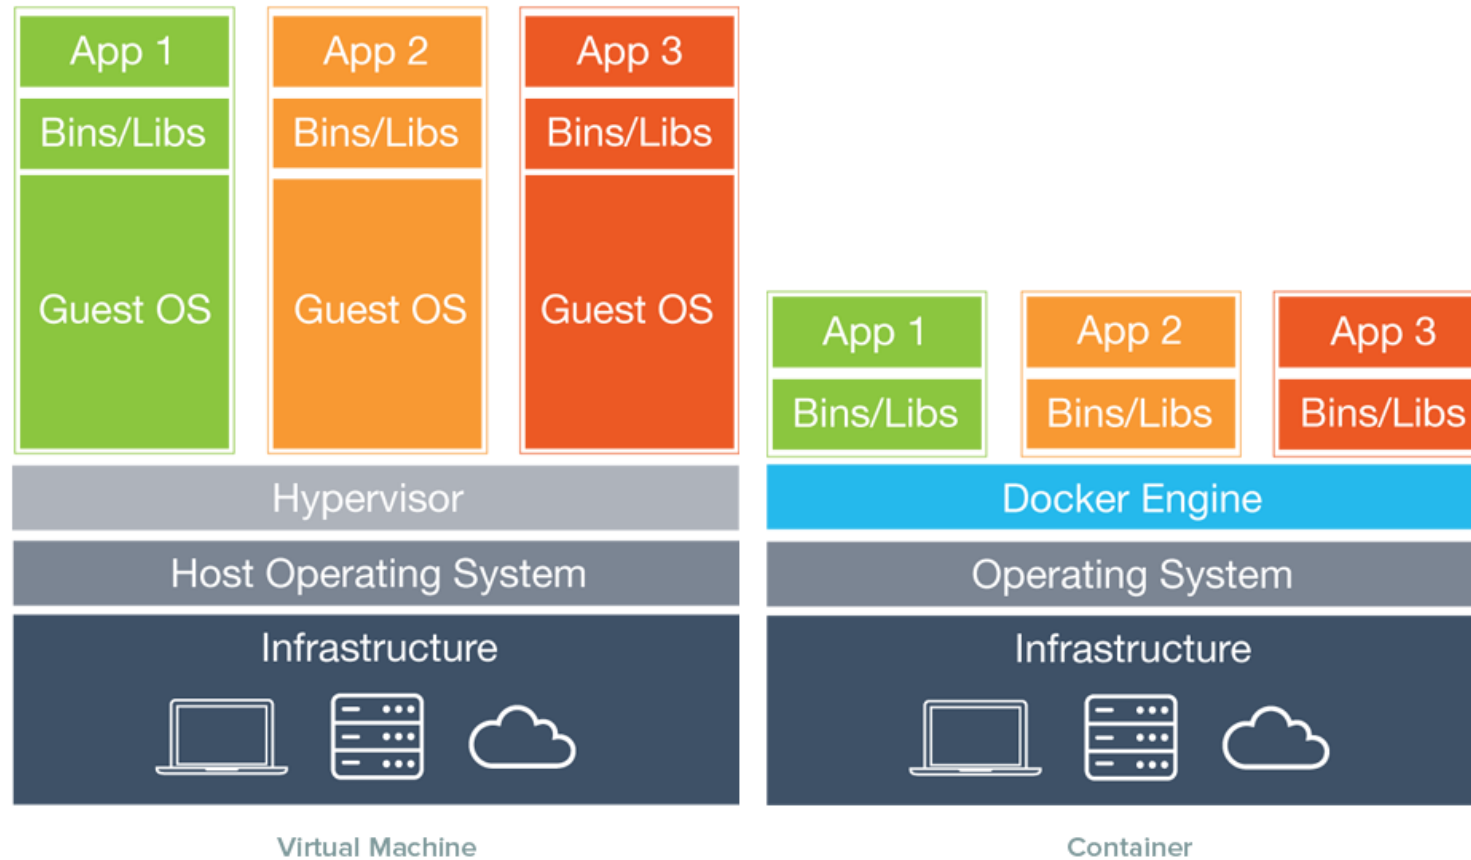
Three core principles

- Kernel Namespaces
 - Linux 2.6 (2009)
- Cgroups
- Chroot

Containers from Scratch – Liz Rice

<https://www.youtube.com/watch?v=MHv6cWjvQjM&t>

Difference to VMs



Difference to VMs

- Decrease virtualization overhead
- No syscall translations
- OS virtualization instead HW virtualization
- Nearly same performance as bare metal
 - Efficiency

Kernel namespaces

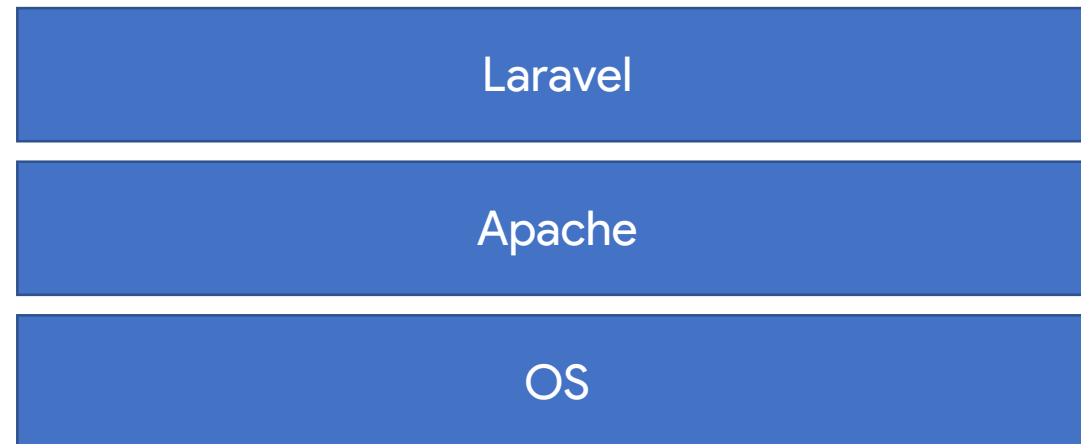
```
[node1] (local) root@192.168.0.48 ~
$ docker ps
CONTAINER ID        IMAGE               COMMAND                  CREATED            STATUS              PORTS
determined_agnesi   bcc1df6c14b2      davidkroell/ytdl       "/usr/bin/env node /..." 19 seconds ago    Up 18 seconds      3000/tcp
determined_agnesi
```

```
[node1] (local) root@192.168.0.48 ~
$ ps -ef | grep node
404 root          0:00 node /usr/src/app/ytdl.js
519 root          0:00 grep node
```


Image

- Package of software
 - with dependencies (libs ...)
 - Includes anything to run
- Application with OS (if needed)
- Storage driver

Example: Laravel Stack




Example: Laravel Stack



Images architecture

- Different Storage Drivers
 - Overlay2 (default)



Exploring Images

- Inspect Images with Docker Toolset
- Use `wagoodman/dive`

Container

- Image which runs on a Docker host (dockerd)
- Network and storage attached
- Consists of an Image
 - Adds container storage layer

Run a container

```
docker run mongo
```

Run a container

```
docker run --name mysql57 -d \  
    --restart=always \  
    -p 127.0.0.1:3306:3306 \  
    -v mysql:/var/run/mysql \  
    -e MYSQL_ROOT_PASSWORD=my-secret-pw \  
mysql:5.7
```

<https://docs.docker.com/engine/reference/run>

Docker Client & Daemon

- JSON API under the hood
 - `/var/run/docker.sock`
- May be called from any HTTP client
- SSH Transport support (v18.09)
- Various integrations
 - Docker2go – <https://davidkroell.github.io/docker2go>

Docker on windows

- Needs HyperV, better use forwarding
- Requires v18.09
- `docker -H ssh://user@host ps`
 - ps on remote server
- Set `DOCKER_HOST` environment variable

Composing an Image

- docker build
- Parse Dockerfile
- Set of statements

```
1  FROM scratch
2  ADD busybox.tar.xz /
3  CMD ["sh"]
```

<https://github.com/docker-library/busybox/blob/master/uclibc/Dockerfile>

Dockerfile

```
FROM bitnami/minideb-extras:stretch-r225
LABEL maintainer "Bitnami <containers@bitnami.com>"

# Install required system packages and dependencies
RUN install_packages ghostscript imagemagick libbz2-1.0 libc6 libcomerr2 libcurl3 libffi6 libfreetype6 libgcc1 libgcrypt2 lib
RUN bitnami-pkg unpack php-7.1.25-20 --checksum b4350f7370f196def8ff56462d7efb4961c15f97a705b96211006a16cec02cac
RUN bitnami-pkg install node-8.14.0-20 --checksum 408efbaecc9a5d5aa93f0e1755a15f5fe29c08d37baaac3900c9f02551d6da2b
RUN bitnami-pkg install laravel-5.7.15-20 --checksum 19e49651d46ceaa1e62c9d03a1d0d0fc369adc6a95867bd88819c1de26e901a2
RUN mkdir /app && chown bitnami:bitnami /app

COPY rootfs /
ENV BITNAMI_APP_NAME="laravel" \
    BITNAMI_IMAGE_VERSION="5.7.15-debian-9-r22" \
    NODE_PATH="/opt/bitnami/node/lib/node_modules" \
    PATH="/opt/bitnami/php/bin:/opt/bitnami/php/sbin:/opt/bitnami/node/bin:/opt/bitnami/laravel/bin:$PATH"

EXPOSE 3000

WORKDIR /app
USER bitnami
ENTRYPOINT [ "/app-entrypoint.sh" ]
CMD [ "php", "artisan", "serve", "--host=0.0.0.0", "--port=3000" ]
```

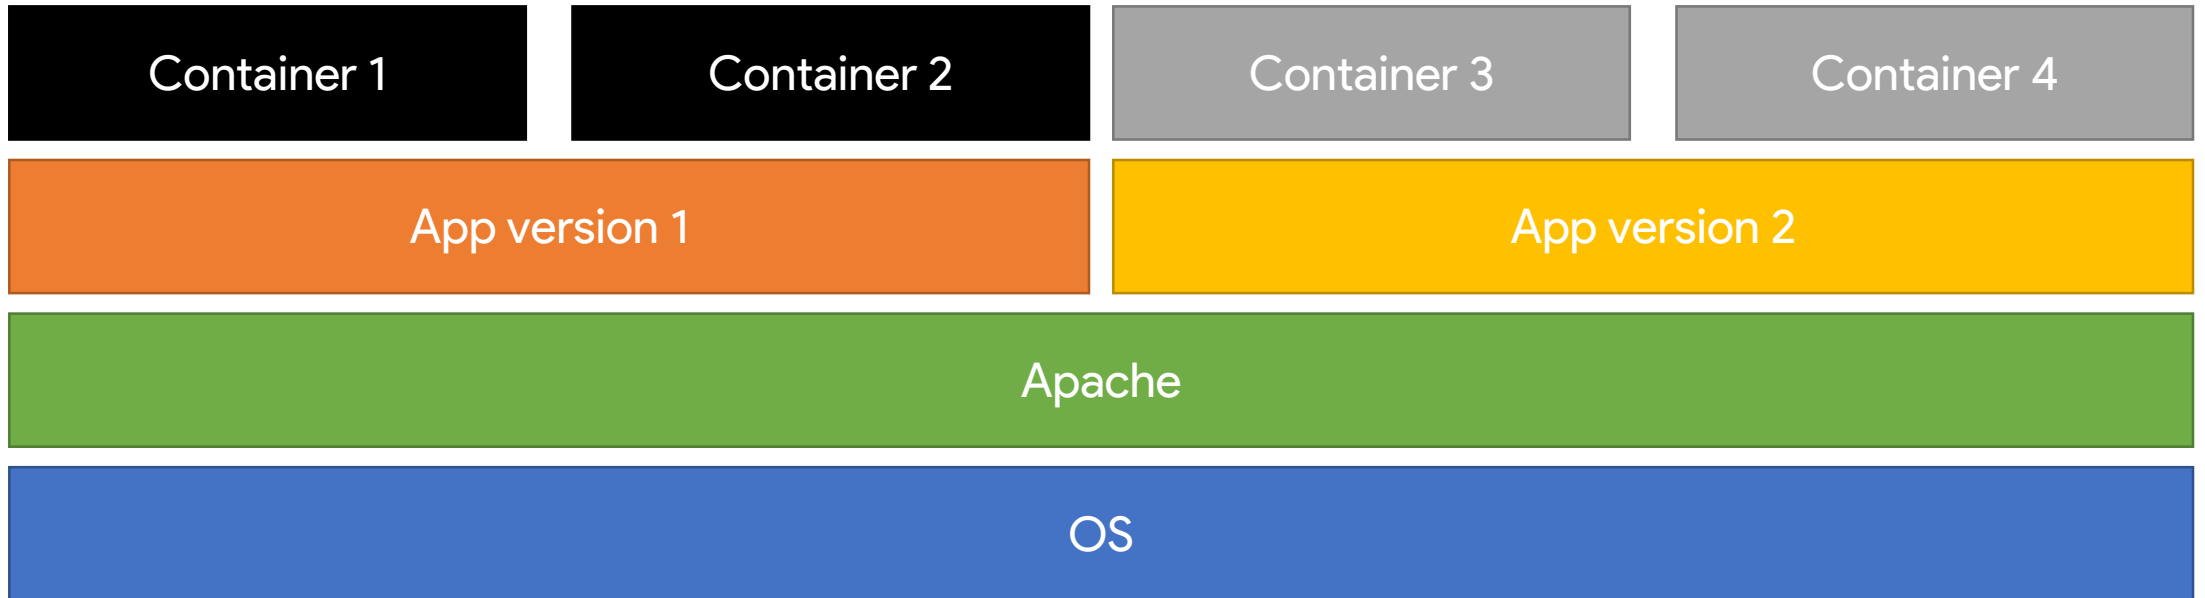
<https://github.com/bitnami/bitnami-docker-laravel/blob/master/5/debian-9/Dockerfile>

Multi-staged build

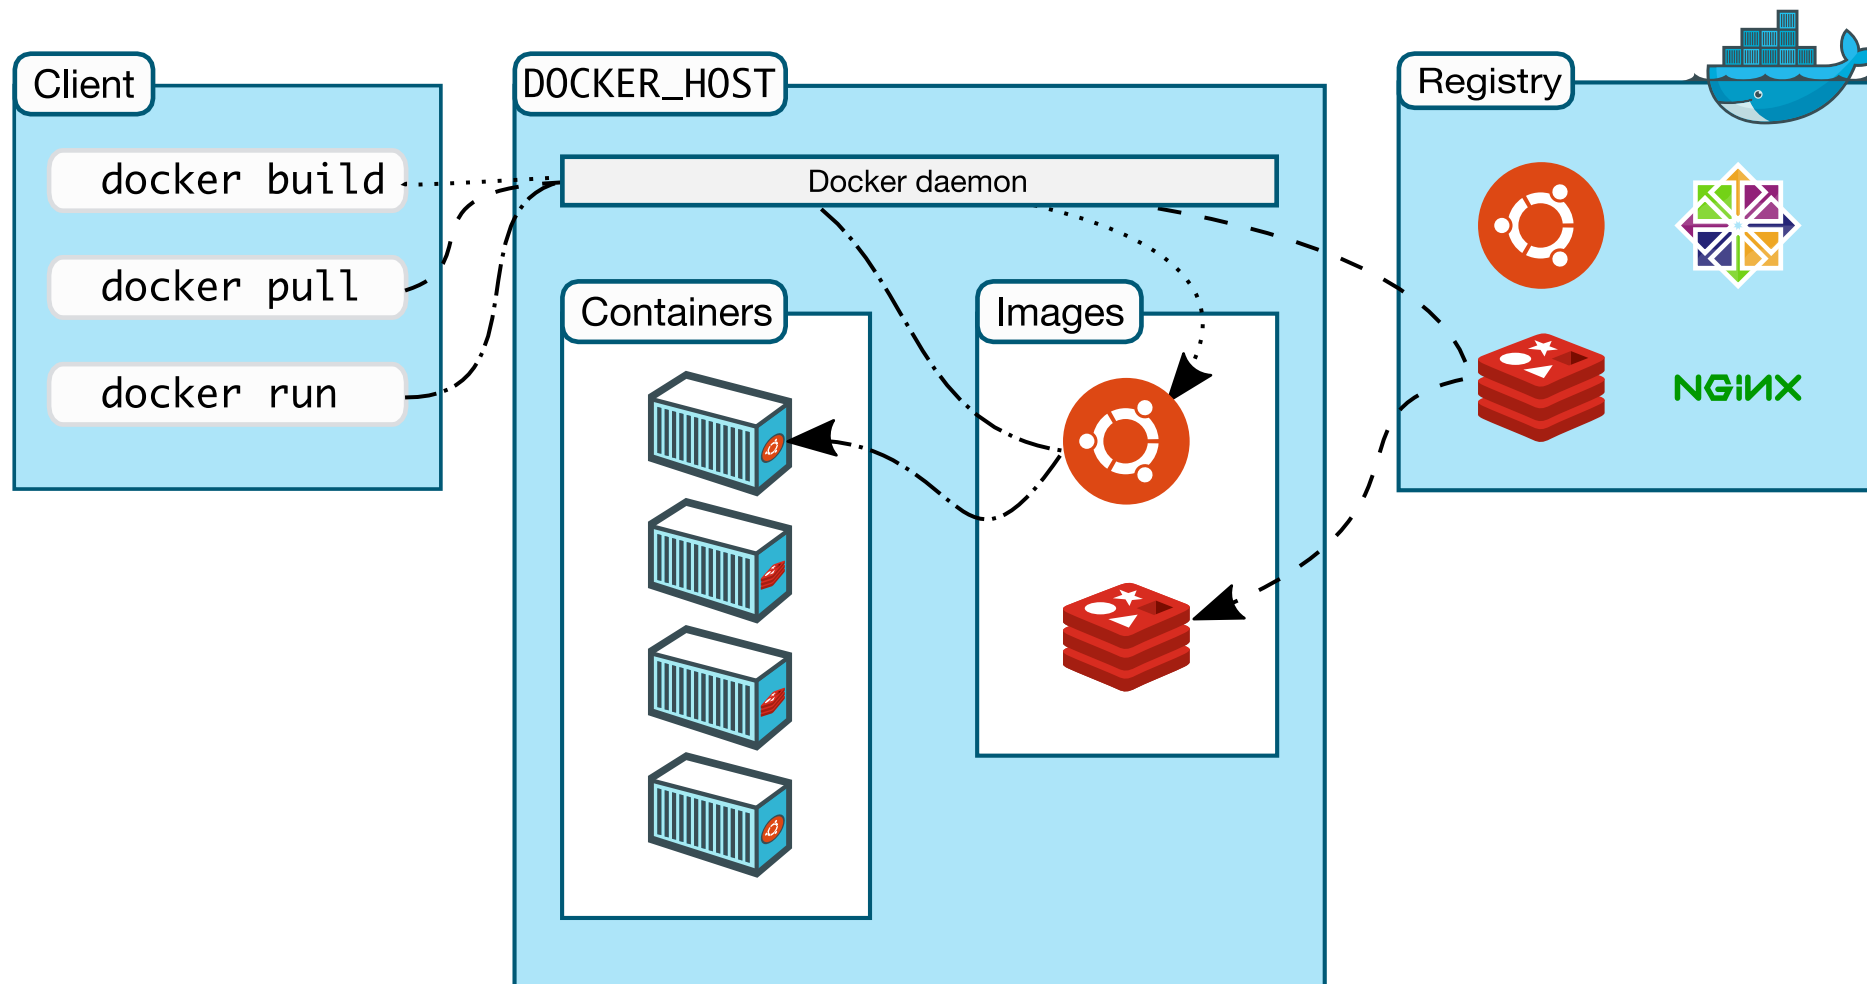
```
1  # base image for building the binary
2  FROM golang:1.11-alpine AS base
3  # $GOPATH is /go
4  COPY . /go/src/github.com/ironarachne/namegen
5  # the output is located in the working directory without fileextension
6  # binary path: /go/main
7  RUN go build /go/src/github.com/ironarachne/namegen/cmd/namegen/main.go
8
9  # add the binary to an empty image
10 FROM scratch
11 # copy from build-image
12 COPY --from=base /go/main /namegen
13 # set namegen as
14 ENTRYPOINT ["/namegen"]
```

<https://github.com/ironarachne/namegen/blob/master/Dockerfile>

Efficiency of Images



Registry Overview



Registry

- Storage for Images
 - Repositories, Tags
- Is delivered as Docker Image
- <https://hub.docker.com> -> official Registry

Registry

```
$ docker run mongo
Unable to find image 'mongo:latest' locally
latest: Pulling from library/mongo
18d680d61657: Pull complete
0addb6fece63: Pull complete
78e58219b215: Pull complete
eb6959a66df2: Pull complete
1bb66a4db707: Pull complete
b18fa018e44b: Pull complete
5d0142bcb0df: Pull complete
db65733de31a: Pull complete
867fe12df2c5: Pull complete
658da3d983f5: Pull complete
fc5dc5157057: Extracting [=====>] 18.94MB/87.01MB
d39407c4ddb7: Download complete
ed6877ef9a9b: Download complete
```

Networking

- Docker implements „virtual“ Networks
- Port mappings
- Inter-container communication
- Address containers with hostnames
 - rrDNS

`docker network ls`

Storage


- Containers should be stateless
- Volumes
- Directory mounts
- Volume plugins for HA storage
 - REX-Ray for S3/GCE/Azure

Deploying Services

- Deployment of dependent containers
- Networking
- Storage
- Environment variables
- Clustering (see Swarm)
- Scaling

docker-compose

- Defines a Service as a whole
- Add-on for Docker
- <https://github.com/docker/compose>



Compose file

```
version: '3'

services:
  web-api:
    image: davidkroell/keylog.rest
    container_name: keylog.rest-api
    restart: always
    environment:
      NODE_ENV: production
      # use this if you are behind a proxy, this is the header, which the proxy sets
      PROXY_HEADER_REAL_IP_KEY: x-real-ip # for jwilder/nginx-proxy
    ports:
      - 3000:3000
```

<https://github.com/davidkroell/keylog.rest/blob/master/docker-compose.yml>

Using a Compose file

- `docker-compose up`
- `docker stack deploy`

Docker Swarm

- Docker-based cluster orchestration
- Multiple Nodes for HA
- Manager/Worker principle (up to 1/1000 ratio)
- Shared nothing
- Fully encrypted overlay networking
- Service discovery
- Rolling updates

Using Docker Swarm

- Enable Swarm mode in dockerd
 - `docker swarm init`
- Add node to cluster
 - `docker swarm join`
- Smallest unit is now a Service, not a container
- Deployment modes and constraints

Deploy to Docker Swarm

- Deployment based on Compose-file
- `docker stack deploy`
- Updates without Downtime
 - Rolling Updates

Scaling your app

- Keeping it stateless makes it simple
- Scale storage => Volume plugins
- Shared-nothing clustering

Monitoring


docker container stats

- Streaming usage API -> JSON
- Cloud-native Monitoring
 - Prometheus
 - cAdvisor
- Limit container resources
 - cgroups

Outlook

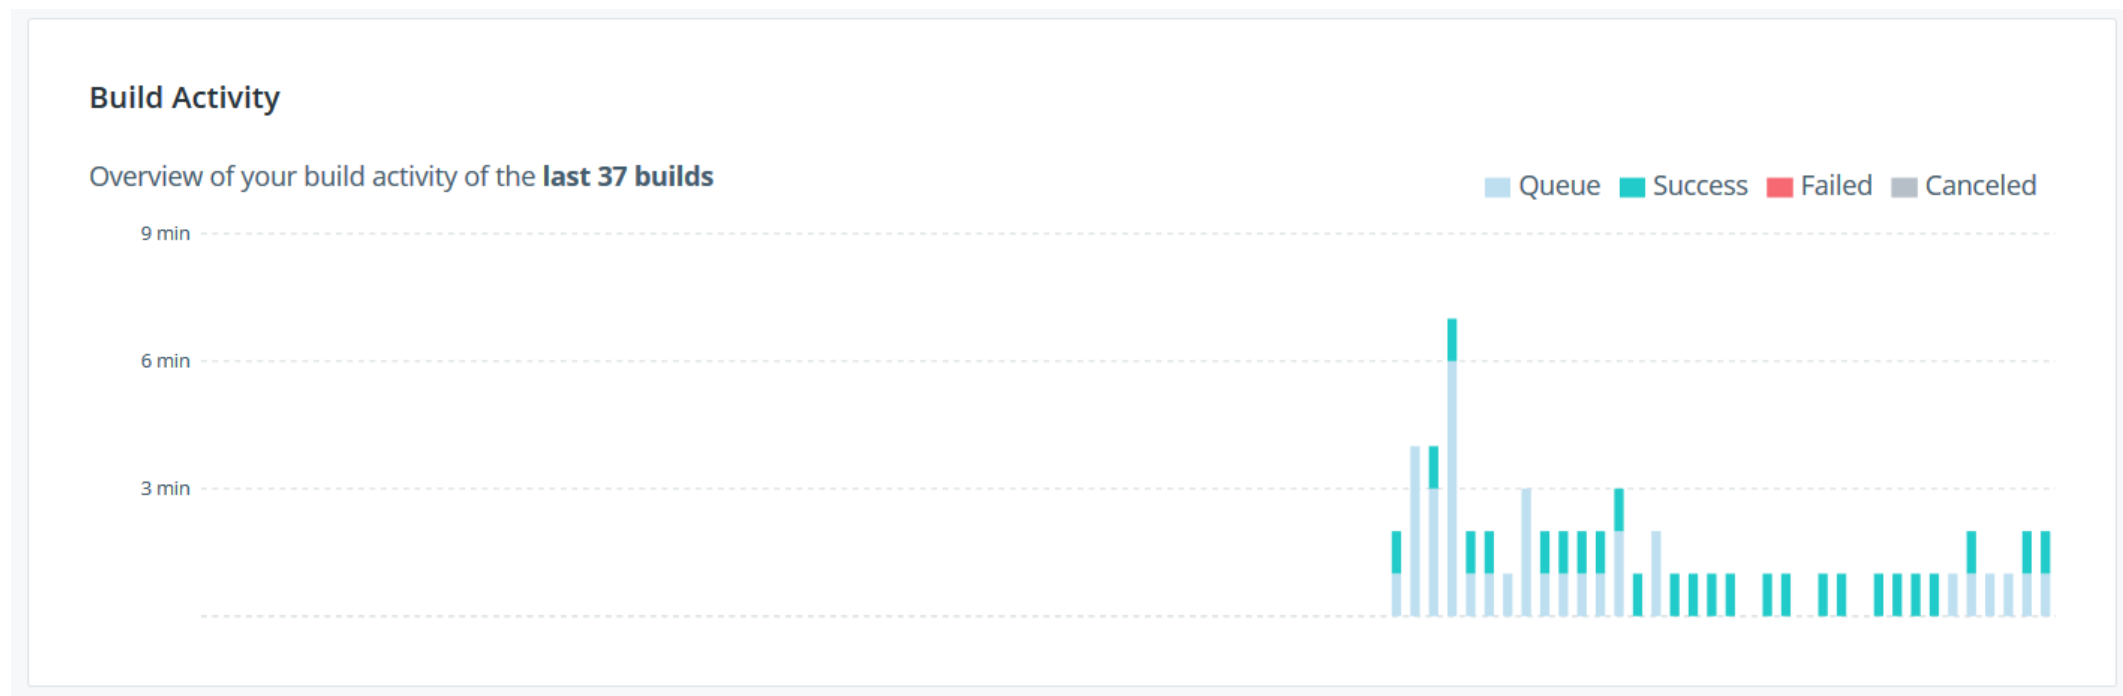
- Kubernetes
 - Global container orchestration platform
- Docker Cloud
- Baremetal cloud
 - CoreOS

Outlook – Kubernetes (K8s)



Outlook – Docker Cloud

– Automated builds



Outlook – Baremetal cloud

- No virtual infrastructure
- Containers
- Orchestrated
 - Docker Swarm
 - K8s or
- CoreOS

Sources

- <https://hackathon.bz.it>
- <https://docs.docker.com>
- <https://www.sdxcentral.com>